

DHANUSH KONDURU

dhanushkonduru@gmail.com | Github: dhanushkonduru | LinkedIn: dhanush-konduru | +91 9676759973 |
LeetCode: DhanushKonduru

SUMMARY

AI/ML and software engineer with three internships behind me, spanning LLM fine-tuning and evaluation, RAG and multi-agent pipelines, and multi-tenant Django backends deployed on AWS and GCP. Cut retrieval token cost by 40% on my own RAG work, won first place at HackX 3.0, and hold two first-named invention disclosures on machine unlearning verification, both cleared as patentable after formal prior-art search.

EDUCATION

Vellore Institute of Technology
Integrated M.Tech in Software Engineering

Vellore, TN
Expected 2027

EXPERIENCE

Software Engineer Intern

Oct 2025 – May 2026

CEEC Global Ventures

Remote

- Built the multi-tenant backend for [Tabzy](#), a cloud ERP/POS product, using custom Django middleware and RBAC so that no company is ever able to see another company's records. Three modules run in production, billing, inventory and orders, with zero cross-tenant leakage across concurrent tenants.
- Shipped an API key system so businesses could connect Tabzy to the tools they already use, and owned the feature end to end from schema design through the production release.
- Sole backend engineer on [BlogSpeed](#), an AI content automation platform. Wrote the auto-scheduler for hands-free publishing, the API key integrations and the content formatting layer, then containerized the whole service and deployed it to GCP through Cloud Build and Cloud Run.
- Tech Stack:** Django, PostgreSQL, Multi-Tenant Architecture, RBAC, Docker, GCP (Cloud Build, Cloud Run)

AI/ML Intern

Jun 2025 – Jul 2025

Centific Global Technologies

Hyderabad, IN (Hybrid)

- Benchmarked five models (GPT, DeepSeek, Phi-2, Phi-3 and Gemma) across several NLP tasks through an automated evaluation harness, comparing latency, throughput and output consistency so model selection came down to measured numbers rather than guesswork about which one felt better.
- Fine-tuned Phi-3 and DeepSeek with LoRA, running hyperparameter sweeps on domain-specific datasets and writing the results up in notebooks that anyone on the team could rerun without asking me first.
- Read through recent work on fine-tuning methods and newer architectures, then turned it into short internal reports the team used when deciding what was worth adopting.
- Tech Stack:** Python, PyTorch, HuggingFace Transformers, scikit-learn, LoRA, OpenAI API, LLM Evaluation

Python Intern

May 2024 – Jun 2024

Aapoon

Hyderabad, IN (Offline)

- Developed a Django and DRF task-management app with OTP login and role-based access for three user types, backed by a normalized schema with validation on every write path, plus REST endpoints that fed the reports on how people were actually using the product from one day to the next.
- Followed that with a college management portal in Django covering student records, attendance marking and separate permission levels for staff, students and the administrative users above them.
- Tech Stack:** Python, Django, Django REST Framework, PostgreSQL

PROJECTS

Financial Document Intelligence RAG | *Python, LangChain, FastAPI, ChromaDB, FAISS, OpenAI, AWS*

- Built a RAG pipeline over SEC 10-K filings and earnings transcripts with LangChain LCEL, pairing BM25 and dense retrieval with a cross-encoder reranker. Faithfulness went from 0.71 to 0.89 on a 50-question RAGAS eval set that also tracked answer relevance and context precision.
- Added contextual compression to filter retrieved chunks before generation, which cut token cost by roughly 40% without losing citation accuracy. Every answer returns its source chunk IDs, so any claim in the output can be traced back to the exact section of the filing that produced it.

Multi-Agent Investment Research Assistant | *Python, LangGraph, OpenAI, SEC EDGAR API, LangSmith, Redis*

- Assembled a LangGraph system with five agents (data gatherer, fundamental analyst, sentiment analyst, risk analyst and report writer) sharing a typed Pydantic state graph, with checkpointing so a run that fails halfway resumes from the last good state instead of starting the whole thing over from scratch.
- Wired the RAG service in as a tool alongside the SEC EDGAR and yfinance APIs, and used LangSmith traces to track down hallucinated tool calls and context overflow failures during agent runs.

Portfolio Construction with Classical ML + MLOps Pipeline | *Python, XGBoost, scikit-learn, MLflow, DVC, Evidently AI*

- Trained an XGBoost cross-sectional return predictor on S&P 500 names using Fama-French style features, validated with expanding-window walk-forward splits, since random k-fold quietly leaks future information into training whenever the underlying data is an ordered financial time series.
- Added Evidently AI to compare live input distributions against the training baseline daily and alert on feature drift, with a FastAPI endpoint serving whichever model is registered as production at the time.

PATENTS (INVENTION DISCLOSURES, PATENTABILITY CLEARED)

Trustworthy Machine Unlearning Through Multi-Prompt Behavioral Assessment for Large Language Models

- Dhanush Konduru, Rajasekhar Babu M, Nirmala M. VIT IPR ID IPR0005034P, patentability search report PAT/2026/1217 issued August 2026, cleared on novelty, inventive step and the Section 3(k) software exclusion. Designs a verification-gated engine that diffs two dataset versions, unlearns the removed records and learns the added ones on one model without retraining, then withholds it until a multi-probe behavioural audit and an adversarial relearning test both pass. Fine-tuning an unlearned model on a few deleted records brought back 56 to 97 percent of the knowledge that was meant to be gone, against 22 percent for a from-scratch retrain, so every method a loss curve would have certified had only hidden the data.

AeroForge: Rapid-Unlearning Digital Twin System for Continuous Fleet Governance and Adversarial Verification

- Dhanush Konduru, Rajasekhar Babu M. VIT IPR ID IPR0005230P, prior-art search report issued August 2026 and assessed patentable on both novelty and inventive step. Carries the verification-gated approach into turbofan fleet prognostics, stripping a departing operator's telemetry out of a remaining-life model in roughly a tenth of the retraining time while the previously approved model stays in service so fleet monitoring never pauses. Adds provenance indexing so an erasure request resolves by lookup rather than similarity search, plus a feasibility check that refuses the request outright when the departing operator's engines turn out to be statistically indistinguishable from the rest of the aircraft on the same monitoring contract.

TECHNICAL SKILLS

Languages: Python, SQL, Java, JavaScript, C++/C, HTML/CSS

AI / Machine Learning: PyTorch, scikit-learn, XGBoost, pandas, NumPy, Jupyter, Deep Learning, NLP, Model Evaluation, Hyperparameter Tuning, Feature Engineering, LoRA Fine-Tuning, HuggingFace Transformers

LLM & Agentic Frameworks: LangChain, LangGraph, LangSmith, OpenAI API, Groq API, RAG Systems, RAGAS Evaluation, Prompt Engineering, Function Calling, Structured Outputs

Backend Development: Django, Django REST Framework, FastAPI, Flask, REST APIs, Multi-Tenant Architecture, RBAC, JWT and OTP Auth, API Key Auth

Databases: PostgreSQL, MongoDB, SQLite, Supabase, Redis, Vector DBs (ChromaDB, Pinecone, FAISS)

MLOps & DevOps: MLflow, DVC, Evidently AI, Model Monitoring, Docker, GitHub Actions, CI/CD, Git

Cloud: AWS (EC2, S3, App Runner), GCP (Cloud Build, Cloud Run)

ACHIEVEMENTS

2025 **1st Place** **HackX 3.0 GraVITas'25, VIT Vellore** (Oracle-sponsored Annual Techno-Management Fest)

2025 **Jury Mention** **ICADT'25 National Level Hackathon, VIT Vellore** (Springer / AdroIT Technologies)

2024 **3rd Place** **National Level Hackathon on Generative AI, NNRG Hyderabad** (IIIC & IIC)

CERTIFICATIONS

Machine Learning Specialization (Stanford / DeepLearning.AI) | **Deep Learning Fundamentals** (NVIDIA) | **Oracle Generative AI** | **Prompt Engineering** (Vanderbilt)